Prospective low-prevalence evaluation

10,000 new URLs • 5.00% phishing • frozen models and thresholds

Condition	Accuracy	Precision	Recall	F1	ROC-AUC	PR-AUC	FPR	FNR
Weighted HYB-MS real-only	97.35%	73.93%	72.60%	73.26%	93.11%	78.35%	1.35%	27.40%
Weighted HYB-MS +10% GAN	97.75%	78.95%	75.00%	76.92%	92.98%	80.91%	1.05%	25.00%

STRICT AUGMENTATION RULE NOT MET

The +10% model did not improve recall, F1 and ROC-AUC while also keeping FPR no higher than real-only. The result is kept as a valid negative or mixed finding.

No training, calibration, checkpoint choice or threshold was changed after this dataset was frozen.

Dataset SHA-256: 6801eb222c49779fb431f4c45e98e85ac984ae8a6bbd65f1c5c29abdc52cb812

What the result means in browser-like traffic

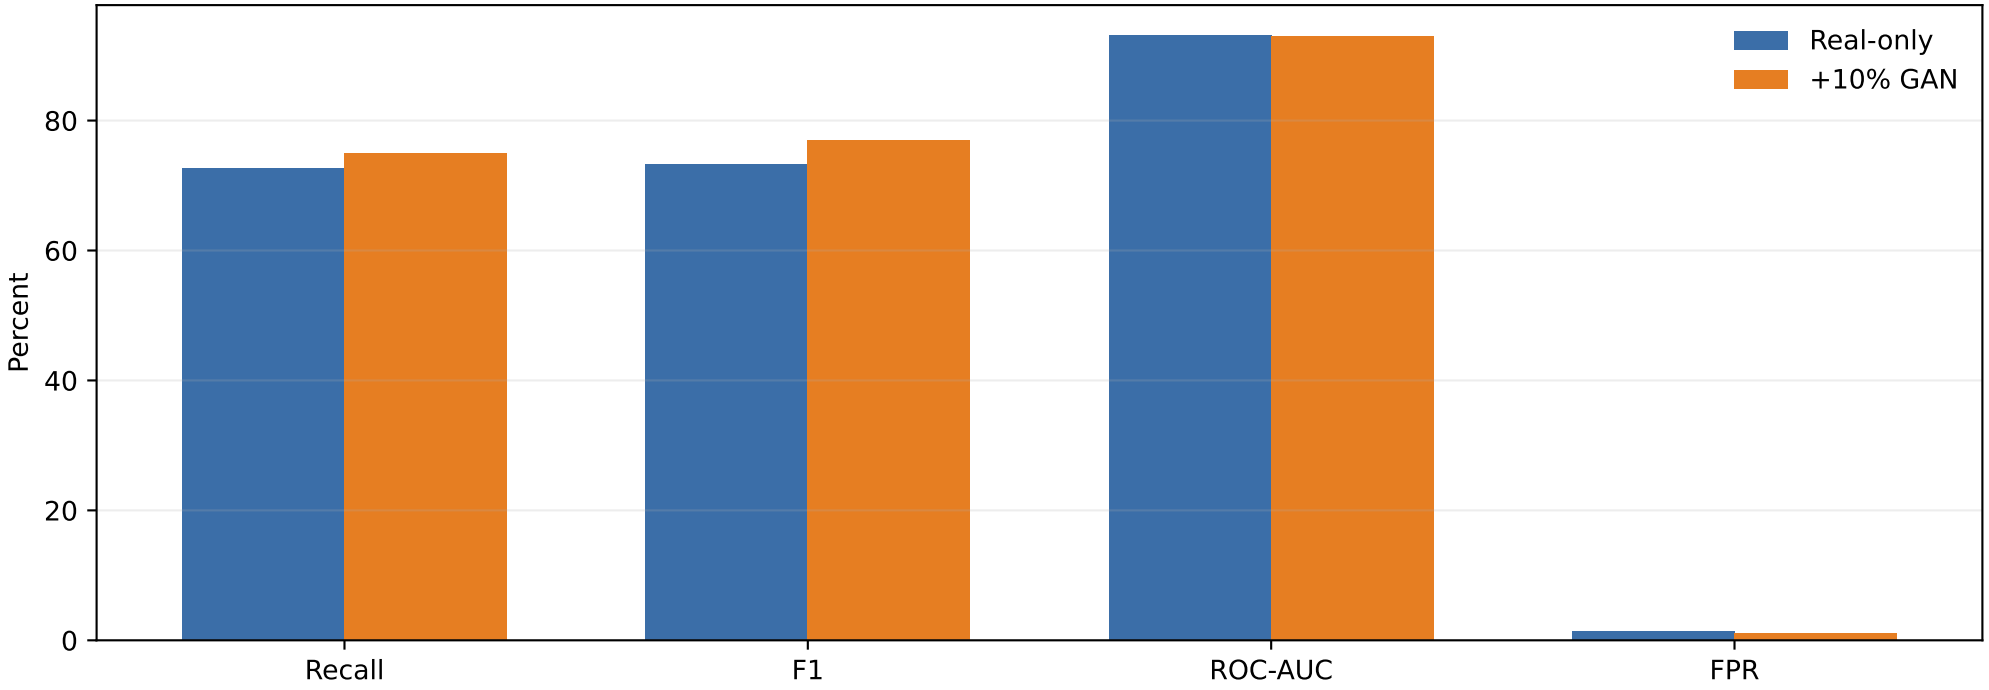

Real-only false alerts per 10,000 visits: 128.0

+10% GAN false alerts per 10,000 visits: 100.0

This is an identity-disjoint, proxy-labelled evaluation. It is useful evidence, but it is still not a live production deployment test.