Prospective low-prevalence evaluation

10,000 new URLs • 5.00% phishing • frozen models and thresholds

Condition	Accuracy	Precision	Recall	F1	ROC-AUC	PR-AUC	FPR	FNR
Weighted HYB-MS real-only	97.07%	73.36%	65.00%	68.93%	83.41%	69.67%	1.24%	35.00%
Weighted HYB-MS +10% GAN	97.47%	79.06%	67.20%	72.65%	82.63%	72.73%	0.94%	32.80%

STRICT AUGMENTATION RULE NOT MET

The +10% model did not improve recall, F1 and ROC-AUC while also keeping FPR no higher than real-only. The result is kept as a valid negative or mixed finding.

No training, calibration, checkpoint choice or threshold was changed after this dataset was frozen.

Dataset SHA-256: fb174d1e7b7aec9b7b780eecb549985420648a52492f1fbad4e7688e6bf752e4

What the result means in browser-like traffic

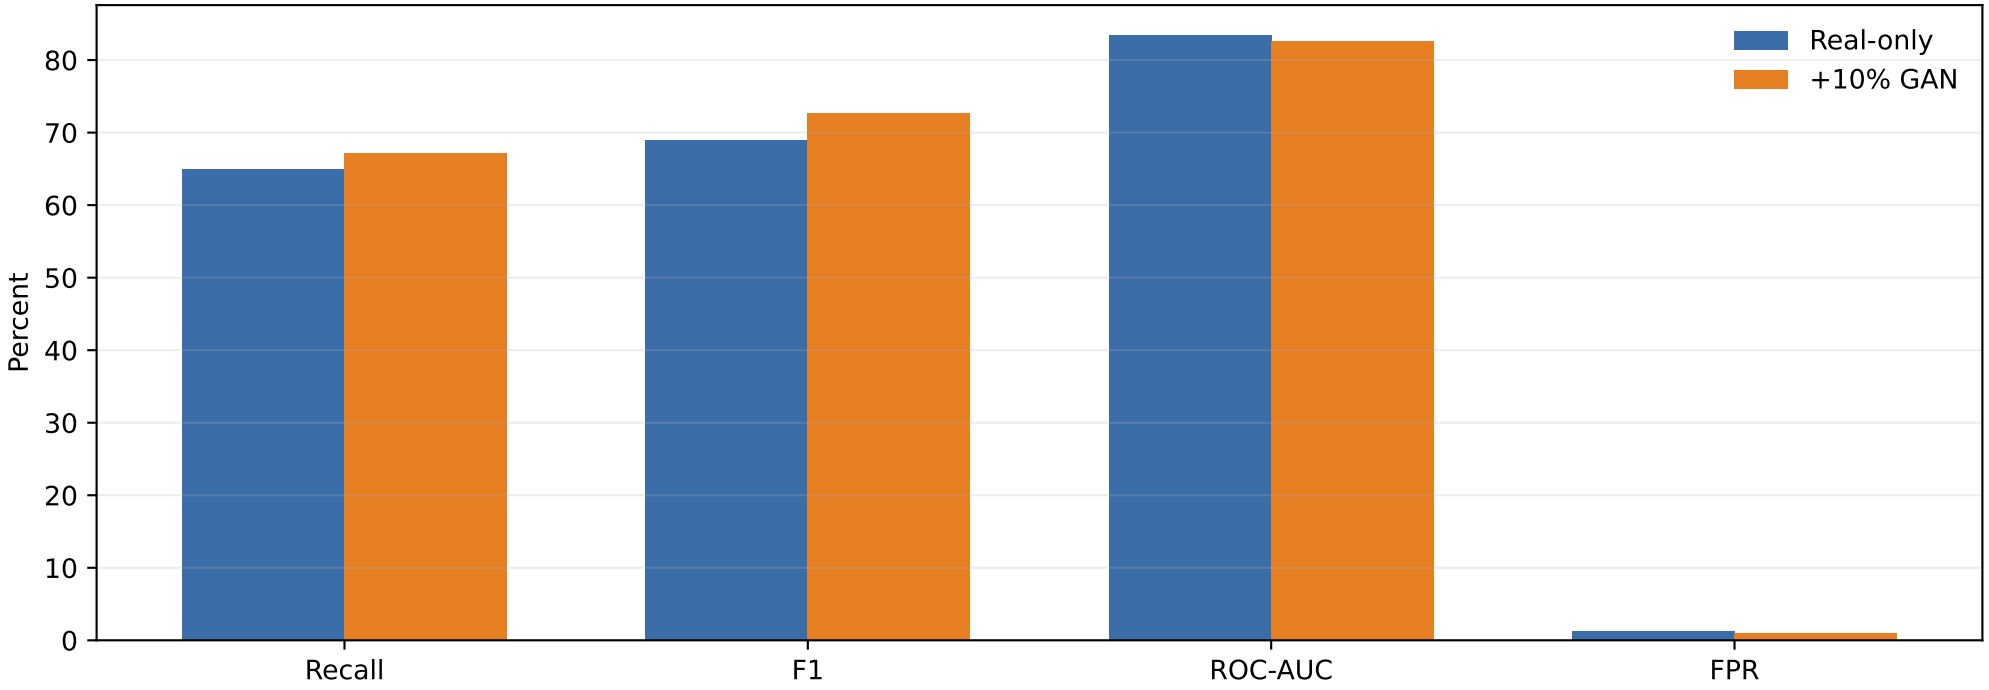

Real-only false alerts per 10,000 visits: 118.0

+10% GAN false alerts per 10,000 visits: 89.0

This is an identity-disjoint, proxy-labelled evaluation. It is useful evidence, but it is still not a live production deployment test.